

3D Scanning & Motion Capture

Exercise - 1

Felix Altenberger, Andrei Burov



Exercise 0 - Outcomes


1. We've received 30% of expected submissions
2. Pitfalls (build output instead of app output)
3. Reach out on Moodle!
4. We hope you tested your working environment now :)
5. We will publish a PDF with matriculation numbers of successful submissions

Exercises – Overview

1. Exercise → Camera Intrinsics, Back-projection, Meshes
2. Exercise → Surface Representations
3. Exercise → Coarse Alignment (Procrustes)
4. Exercise → Optimization
5. Exercise → Object Alignment, ICP

Exercise 1

1. Exercise → Camera Intrinsics, Back-projection, Meshes



TUM-RGB-D SLAM Dataset

- <https://vision.in.tum.de/data/datasets/rgbd-dataset>
- 39 sequences
- Recorded using Kinect v.1
 - Structured light
 - Calibrated
 - Aligned depth and color maps
- Camera trajectory




TUM-RGB-D SLAM Dataset




3D Scanning & Motion Capture
Thies, Dai, Altenberger, Burov

Scene: "fr1/xyz"

Kinect v.1 – Depth and Color Information



Kinect v.1 – Structured Light



Tasks

1. Project dependencies & CMake configuration
2. Back-Projection
 - Use the given intrinsics, extrinsics and the camera trajectory to project the camera observation back to world space
 - Assign the color to the back-projected points
3. Write a 3D mesh
 - Write an OFF file containing the back-projected position and color information
 - Make use of the grid structure of the observation to perform the triangulation

Task 0) Project dependencies

- Eigen <http://eigen.tuxfamily.org>
 - Headers-only
 - Linear Algebra library
 - Matrix, Vector, Solvers, ...
 - TIP: Do not use C++'s `auto`
- FreeImage <http://freeimage.sourceforge.net/>
 - Support for many image formats
 - Windows: We provide a pre-compiled binary
 - Linux: `$ sudo apt-get install libfreeimage3 libfreeimage-dev`

Task 1) Back-Projection

- Use depth map, camera intrinsics and trajectory to project points from 2D \rightarrow 3D.



1 float / pixel (z)




4 chars / pixel (R, G, B, A)



Point in 3D / world space


How are images synthesized?

- Given a real-world/CG scene and a camera, we want to project the 3D points in the scene to 2D pixel positions in the image





World Space

- Every point in the scene has its (X, Y, Z) coordinates




World -> Camera Space

- We place a camera at the (X, Y, Z) position in world space
- The pose/orientation is given by the right/up/look vectors




World -> Camera Space


- To transform points from world space to camera space, we need to multiply them by the camera extrinsic matrix




$$\text{Extrinsic} = \text{Pose}^{-1}$$




Camera Extrinsic matrix



Pinhole camera model




Pinhole camera model



Camera Intrinsic Matrix

$$\begin{pmatrix} x' \\ y' \\ z' \end{pmatrix} = K \cdot \begin{pmatrix} X_c \\ Y_c \\ Z_c \end{pmatrix}$$

$$\begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} x'/z' \\ y'/z' \end{pmatrix}$$



Intrinsic matrix

$f := \text{focal length} = 4.1mm$

$W := \text{sensor width} = 4.54mm$

$H := \text{sensor height} = 3.42mm$

$w := \text{image width} = 640$

$h := \text{image height} = 480$

$c_x := \text{image center x} = 320$

$c_y := \text{image center y} = 240$


Resulting intrinsic matrix :

$$\begin{bmatrix} \frac{f \cdot w}{W} & 0 & c_x \\ 0 & \frac{f \cdot h}{H} & c_y \\ 0 & 0 & 1 \end{bmatrix}$$


Perspective Projection

$$\begin{pmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} X_c \\ Y_c \\ Z_c \end{pmatrix} = \begin{pmatrix} x' \\ y' \\ z' \end{pmatrix} \quad \xrightarrow{\text{Dehomogenization}} \quad \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} x'/z' \\ y'/z' \end{pmatrix}$$

Projection Pipeline



Projection Pipeline



More Than One Camera

- RGB-D Sensor like the Kinect (or your phone) have multiple cameras
- This raises the question: Which camera does the extrinsic matrix correspond to?

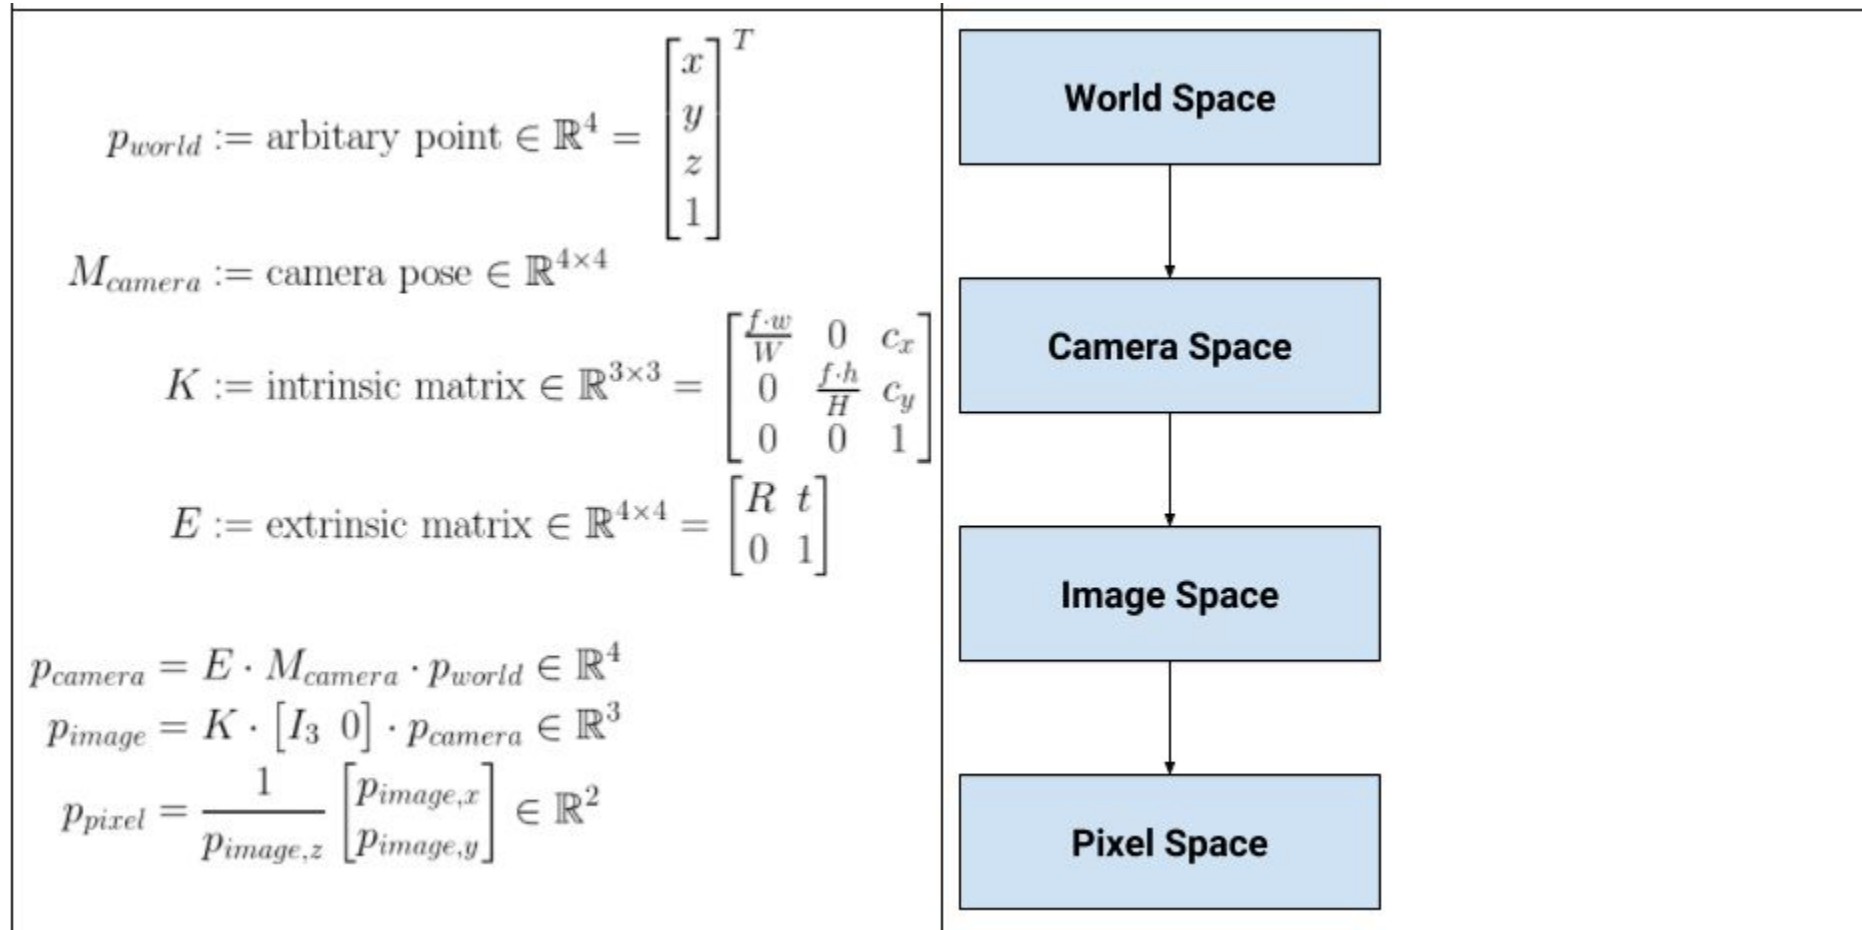


More Than One Camera

- Pick a point on the camera as the reference
- Each camera gets an extrinsic matrix that describes its pose relative to the reference point



Projection Pipeline (from World to Pixels)




Task 2) Mesh Output

- Write OFF file
 - Simple text-based format
 - Vertices/Points:
 - Position
 - Color
 - Faces
 - Indices of vertices

```
1  COFF
2  # numVertices numFaces numEdges
3  4 2 0
4  # list of vertices
5  # X Y Z R G B A
6  0.0 1.0 0.0 255 255 255 255
7  0.0 0.0 0.0 255 255 255 255
8  1.0 0.0 0.0 255 255 255 255
9  1.0 1.0 0.0 255 255 255 255
10 # list of faces
11 # nVerticesPerFace idx0 idx1 idx2 ...
12 3 0 1 2
13 3 0 2 3
```

Task 2) Mesh Structure



Ensure consistent
orientation of the triangles!
(Usually counter-clockwise)

Example:

First triangle: 0-5-1

Second triangle: 5-6-1

Task 3) Submit your solution to Moodle

- Upload your main.cpp and a screenshot from MeshLab to Moodle
- If you worked in a group
 - **Each member** of the group should upload the solution
 - List all team members **names** and **matriculation numbers** in a separate **team_members.txt** file and upload it with your solution

See you next time!